

What is Amazon Athena?

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard [SQL](#). With a few actions in the AWS Management Console, you can point Athena at your data stored in Amazon S3 and begin using standard SQL to run ad-hoc queries and get results in seconds.

Amazon Athena also makes it easy to interactively run data analytics using Apache Spark without having to plan for, configure, or manage resources. When you run Apache Spark applications on Athena, you submit Spark code for processing and receive the results directly.

Athena SQL and Apache Spark on Amazon Athena are serverless, so there is no infrastructure to set up or manage, and you pay only for the queries you run. Athena scales automatically—running queries in parallel—so results are fast, even with large datasets and complex queries.

Athena helps you analyze unstructured, semi-structured, and structured data stored in Amazon S3. Examples include CSV, JSON, or columnar data formats such as Apache Parquet and Apache ORC. You can use Athena to run ad-hoc queries using ANSI SQL, without the need to aggregate or load the data into Athena.

Athena integrates with Amazon QuickSight for easy data visualization. You can use Athena to generate reports or to explore data with business intelligence tools or SQL clients connected with a JDBC or an ODBC driver.

Athena integrates with the AWS Glue Data Catalog, which offers a persistent metadata store for your data in Amazon S3. This allows you to create tables and query data in Athena based on a central metadata store available throughout your Amazon Web Services account and integrated with the ETL and data discovery features of AWS Glue.

Amazon Athena makes it easy to run interactive queries against data directly in Amazon S3 without having to format data or manage infrastructure. For example, Athena is useful if you want to run a quick query on web logs to troubleshoot a performance issue on your site. With Athena, you can get started fast: you just define a table for your data and start querying using standard SQL.

You should use Amazon Athena if you want to run interactive ad hoc SQL queries against data on Amazon S3, without having to manage any infrastructure or clusters. Amazon Athena provides the easiest way to run ad hoc queries for data in Amazon S3 without the need to setup or manage any servers.

AWS Athena is a powerful serverless query service provided by AWS for analyzing the data directly in Amazon S3 using standard SQL. It facilitates features like high scalability, cost-effectiveness, easy-to-use platform for running complex queries without the need for extensive infrastructure setup. In this article we will discuss on what is aws athena, its architecture, benefits, limitations, advantages, disadvantages and how it difference from Amazon Redshift, Amazon Glue and Microsoft SQL server effectively.


AWS Athena is a serverless interactive query service that enables normal SQL data analysis in Amazon S3. Athena is based on Presto, a distributed SQL query engine, and it can query data in Amazon S3 fast using conventional SQL syntax. There is no infrastructure to handle with Athena, so

you can focus on analyzing data at scale. To have more idea of AWS Athena, let us understand the architecture first.

AWS Athena Architecture

Apache Presto, an open-source distributed SQL query engine, serves as the foundation for Athena. When a query is submitted by a user, Athena generates a query plan and sends it to Presto for execution. Presto then distributes the query over numerous cluster nodes for parallel processing. The results are subsequently compiled and presented to the user. Athena stores table and partition metadata in a controlled Hive metastore.

When a query is run, Athena gets the metadata from the metastore to establish the data's location and format. Athena also interfaces with AWS Glue, a fully managed extract, transform, and load (ETL) service, allowing customers to create and manage data catalogs and ETL processes. Furthermore, we will go through the various components of AWS Athena.



Amazon S3: Athena searches data stored in Amazon S3, an object storage service that is highly durable, highly accessible, and infinitely scalable.

Amazon Glue: Athena leverages AWS Glue, a fully managed extract, transform, and load (ETL) service, to catalog and query the data stored in S3.

Apache Presto: Apache Presto is Athena's distributed SQL query engine. Presto is well-suited for querying data stored in distributed systems and can handle queries that require data from numerous sources to be joined.

Amazon CloudWatch: Athena interacts with Amazon CloudWatch, a monitoring service that offers metrics and logs for all of your AWS account's resources. CloudWatch may be used to track the performance of your Athena queries and create alerts for specific query patterns.

Amazon VPC: Athena supports performing queries within an Amazon Virtual Private Cloud (VPC), which allows you to isolate your data and limit access to it using Amazon VPC security groups and network ACLs.

Encryption: Athena supports S3 server-side encryption with Amazon S3-managed keys (SSE-S3) or AWS Key Management Service-managed keys (SSE-KMS), as well as SSL/TLS encryption of data in transit.

What are the benefits of using Amazon Athena?

The following are the benefits of using Amazon Athena:

- **Serverless:** For using Amazon Athena no infrastructure management is needed as it automatically handles the scaling, patching and configuration.
- **Cost-Effective:** It only charges for the queries you run with no upfront costs or resource provisioning.
- **Ease of Use:** it query the data directly from the Amazon S3 using standard SQL that is accessible to users familiar with SQL.
- **No infrastructure setup:** Athena is a serverless service that eliminates the need for users to set up and manage infrastructure, making data querying easier and faster.
- **Cost-effective:** Athena charges customers solely for the quantity of data scanned by their searches, making it an affordable solution for ad hoc and exploratory queries.
- **Scalability:** Athena is a fully-managed service that can automatically scale to accommodate massive amounts of data and queries.
- **SQL support:** Since Athena supports ANSI SQL, users can query data in S3 using their existing SQL knowledge and tools.

What are some Amazon Athena Limitations?

The following are the some limitations of Amazon Athena:

- **Comple Queries:** It performance mya be degraded with highly complex queries or ery large datasets that requires multiple joins and aggregations.
- **Cold Start Latency:** Its initial query exection may experience some delay due to its cold start latency, especially for the infrequently accessed data.
- **Limited Data Manipulation:** Athena is primarily useful for querying and doesn't support any data modification operations like INSERT,, UPDATE, or DELETE.

- **Restricted query performance:** The volume of data scanned and the intricacy of the query can limit Athena's speed, resulting in lengthier query times.
- **No real-time querying:** Because Athena is intended for batch processing, it may not be ideal for real-time querying.
- **Limited data types:** In comparison to other database systems, Athena only supports a restricted selection of data types.

Features of AWS Athena

The following are the features of AWS Athena:

- **Serverless architecture:** Athena is a fully-managed service that does not require any infrastructure setup, management, or scaling.
- **Standard SQL support:** Since the Athena supports of ANSI SQL, users can easily query the data in S3 through using their existing SQL knowledge and tools.
- **Connection with the AWS ecosystem:** Athena interfaces with other AWS services such as Amazon S3, AWS Glue, and [AWS Lambda](#), enabling customers to import and convert data from a variety of sources.
- **Cost-effective pricing model:** Athena's pricing approach is cost-effective since it costs customers based on the amount of data scanned by their queries, making it ideal for ad-hoc and exploratory queries.
- **Integration with BI tools:** Athena provides connectivity with major business intelligence tools like as Tableau, [Power BI](#), and Amazon QuickSight, allowing users to build visualizations and reports.

Features	Amazon Athena	AWS Redshift	Microsoft SQL Server	AWS Glue
Service Type	It is servless interactive query service	It is fully managed by data warehouse.	It is relational database management system	is is serverless data integration service.
Primary Use Case	It performs adhoc querying on Amazon S3	It facilitates in Data warehousing and OLAP	It facilitates with transactional and analytical processing	It facilitates with ETL and data cataloging.

Features	Amazon Athena	AWS Redshift	Microsoft SQL Server	AWS Glue
Pricing Model	Pay per query based on data scanned (\$5/TB)	Pay per node/hour and additional storage costs	Licensing costs and pay-per-usage for cloud	Pay per usage (job runs, data catalog storage)
Data Storage	Amazon S3	Redshift managed storage, integrates with S3	Local or cloud storage, depends on setup	Amazon S3 and other data sources
Performance	Optimized for quick queries on large datasets	High performance for complex queries and large datasets	High performance for transactional and analytical workloads	Optimized for ETL operations and data transformation
Maintenance	Fully managed, no maintenance required	Managed service, but requires some administration	Requires regular maintenance and updates	Fully managed, minimal maintenance required
Data Formats Supported	JSON, CSV, Parquet, ORC, Avro	JSON, CSV, Parquet, ORC, Avro, and more	Traditional RDBMS formats	JSON, CSV, Parquet, ORC, Avro, and more

What is AWS Data Exchange?

AWS Data Exchange is a service that helps AWS customers easily share and manage data entitlements from other organizations at scale.

As a data receiver, you can track and manage all of your data grants and AWS Marketplace data subscriptions in one place. When you have access to an AWS Data Exchange data set, you can use compatible AWS or partner analytics and machine learning to extract insights from it.

For data senders, AWS Data Exchange eliminates the need to build and maintain any data delivery and entitlement infrastructure. Anyone with an AWS account can create and send data grants to data receivers. To sell your data as a product in AWS Marketplace, make sure that you follow the guidelines to determine eligibility.

What is a data grant in AWS Data Exchange

A data grant is the unit of exchange in AWS Data Exchange that is created by a data sender in order to grant a data receiver access to a data set. When a data sender creates a data grant, a grant request is sent to the data receiver's AWS account. A data receiver accepts the data grant to gain access to the underlying data.

A grant has the following parts:

- **Data set** – A data set in AWS Data Exchange is a resource curated by the sender. It contains the data assets a receiver will gain access to after accepting a data grant. AWS Data Exchange supports five types of data sets: Files, API, Amazon Redshift, Amazon S3, and AWS Lake Formation (Preview).
- **Data grant details** – This information includes a name and description of the data grant that will be visible to data receivers.
- **Recipient access details** – This information includes the receiver's AWS account ID and specifies how long the receiver should have access to the data.

What is an AWS Marketplace data product?

A product is the unit of exchange in AWS Marketplace that is published by a provider and made available for use to subscribers. A data product is a product that includes AWS Data Exchange data sets. When a data provider publishes a data product, that product is listed in the AWS Marketplace product catalog after being reviewed by

AWS against our guidelines and terms and conditions. Each product published is uniquely identified by its product ID.

A data product has the following parts:

- **Product details** – This information includes name, descriptions (both short and long), data samples, a logo image, and support contact information. Providers complete the product details.
- **Product offers** – Offers define the terms that subscribers are agreeing to when they subscribe to a product. To make a product available in the public AWS Marketplace Catalog, providers must define a public offer. This offer includes prices and durations, data subscription agreement, refund policy, and the option to create custom offers.
- **Data sets** – A product can contain one or more data sets. A data set in AWS Data Exchange is a resource curated by the data provider and contains the data assets a receiver will gain access to after accepting a data grant. AWS Data Exchange supports five types of data sets: Files, API, Amazon Redshift, Amazon S3, and AWS Lake Formation (Preview).

Supported data sets

AWS Marketplace takes a responsible approach to facilitating data transactions by promoting transparency through use of the service. AWS Marketplace reviews permitted data types, restricting products that are not permitted. Providers are limited to distributing data sets that meet the legal eligibility requirements set forth in the Terms and Conditions for AWS Marketplace Sellers.

Accessing AWS Data Exchange

Data receivers

As a data receiver, you can view all of your current, pending, and expired data grants from the AWS Data Exchange console.

You can also discover and subscribe to new third-party data sets available through AWS Data Exchange from the [AWS Marketplace catalog](#).

Data senders and providers

As a data sender or provider, you can access AWS Data Exchange through the following options:

- Directly through the [AWS Data Exchange console](#) (**Publish data**)
- Data providers with data products available in AWS Marketplace can access programmatically using the following APIs:
 - **AWS Data Exchange API** – Use the API operations to create, view, update, and delete data sets and revisions. You can also use these API operations to import and export assets to and from those revisions.
 - **AWS Marketplace Catalog API** – Use the API operations to view and update data products published to AWS Marketplace.

Supported Regions

AWS Data Exchange data grants, subscriptions, data sets, revisions, and assets are Region resources that can be managed programmatically or through the AWS Data Exchange console in supported Regions. Data products published to AWS Marketplace are available in a single product catalog. Subscribers can see the same catalog regardless of which supported AWS Region they are using.

The following regions are supported:

- US East (N. Virginia)
- US East (Ohio)
- US West (Oregon)
- US West (N. California)
- EU (Ireland)
- EU (Frankfurt)
- EU (London)
- Asia Pacific (Singapore)
- Asia Pacific (Tokyo)
- Asia Pacific (Sydney)
- Asia Pacific (Seoul)

The following services are related to AWS Data Exchange:

- **Amazon S3** – AWS Data Exchange allows providers to import and store data files from their Amazon S3 buckets. Data recipients can export these files to Amazon S3 programmatically. AWS Data Exchange also enables recipients to directly access and use providers' Amazon S3 buckets.
- **Amazon API Gateway** – Another supported asset type for data sets is APIs. Data recipients can call the API programmatically, call the API from the AWS Data Exchange console, or download the OpenAPI specification file.

- **Amazon Redshift** – AWS Data Exchange supports Amazon Redshift data sets. Data recipients can get read-only access to query the data in Amazon Redshift without extracting, transforming, and loading data.
- **AWS Marketplace** – AWS Data Exchange allows data sets to be published as products in AWS Marketplace. AWS Data Exchange data providers must be registered as AWS Marketplace sellers, and can use the AWS Marketplace Management Portal or the AWS Marketplace Catalog API.
- **AWS Lake Formation** – AWS Data Exchange supports AWS Lake Formation data permission data sets (Preview). Data recipients get access to data stored in a data provider's AWS Lake Formation data lake and can query, transform, and share access to this data from their own AWS Lake Formation data set.

Amazon EMR

Amazon EMR makes it simple and cost effective to run highly distributed processing frameworks such as Hadoop, Spark, and Presto when compared to on-premises deployments. Amazon EMR is flexible – you can run custom applications and code, and define specific compute, memory, storage, and application parameters to optimize your analytic requirements.

In addition to running SQL queries, Amazon EMR can run a wide variety of scale-out data processing tasks for applications such as machine learning, graph analytics, data transformation, streaming data, and virtually anything you can code. You should use Amazon EMR if you use custom code to process and analyze extremely large datasets with the latest big data processing frameworks such as Spark, Hadoop, Presto, or Hbase. Amazon EMR gives you full control over the configuration of your clusters and the software installed on them.

You can use Amazon Athena to query data that you process using Amazon EMR. Amazon Athena supports many of the same data formats as Amazon EMR. Athena's data catalog is Hive metastore compatible. If you use EMR and already have a Hive metastore, you can run your DDL statements on Amazon Athena and query your data immediately without affecting your Amazon EMR jobs.

Amazon Elastic MapReduce is an important cloud-based platform service that is designed for the effective scaling and processing of large-volume datasets. Its platform facilitates the users in quickly and easily setting up the cluster with Amazon


EC2 Instances that are already pre-configured with big data frameworks. In this article, you will explore the easy setup and administration of EMR clusters in AWS.

[Amazon EMR](#) (Elastic Map Reduce) is an AWS-based platform service that processes large-volume datasets using shared computing frameworks such as [Apache Hadoop](#) and [Apache Spark](#). It facilitates the users in quickly setting up, configuring, and scaling virtual server clusters for analyzing and processing vast amounts of data efficiently.

How Does Amazon EMR Work?

Amazon EMR functionalities simplify the complex processing of large datasets over the cloud. Users can create the clusters and can be utilized with elastic nature of [Amazon EC2](#) instances. The natures of Amazon EC2 instances are configured with pre existing frameworks like Apache Hadoop and Apache Spark. By distributing the processing jobs across the several nodes these clusters effectively handle and guarantee the parallel executions with faster outcomes. It provides scalability by automatically adjusting the cluster size in accordance to workload needs. It optimizes the data storages on integrating with other [AWS services](#) making things easier. Users can find the things easily rather than going for complicated detailing of infrastructure and administration. It provides a simplified approach for [big data analytics](#).

Amazon EMR



Amazon EMR Architecture

Amazon EMR (Elastic MapReduce) architecture is designed for efficient big data processing using a distributed computing framework.

1. **Clusters:** Consist of a master node (manages the cluster), core nodes (process data and store data in [HDFS](#)), and optional task nodes (handle additional processing).
2. **Hadoop Ecosystem:** Utilizes tools like [Apache Spark](#), [HBase](#), and Hive, pre-configured and optimized for big data analytics.
3. **AWS Integration:** Seamlessly integrates with [AWS services](#) like S3 (storage), [IAM](#) (security), [CloudWatch](#) (monitoring), and [Amazon VPC](#) (network isolation), enhancing functionality and security.

Features of Amazon EMR

The following are the popular features of Amazon EMR:

- **Integration:** It support integration with other AWS services that enhances the efficiency in data processing, making connections with Amazon S3 possible facilitating efficiency in workflow.
- **Salability:** Amazon EMR providing scaling and handling of workloads dynamically. It support automatic adjustments in sizing of the cluster and optimizing the performance and minimizing costs.
- **Ease Of Use:** Amazon EMR makes the deployments of big data easier by offering pre-configured environments for Apache Hadoop and Apache spark. Setuping and maintaining of clusters will be easier for users without requirement of complex setups on this Amazon ECR.
- **Cost Management:** EMR facilitates with cost optimization through letting users to pay only for the resources during the processing of big data making analytics more affordable. Spot instances and Reserved Instances further minimizes the costs.
- **Security:** EMR provides strong security features such as Data encryption, IAM roles and fine-grained access controls. It ensures data protection through the pipeline processing.

Deployment Options of Amazon EMR

Amazon EMR offers many different deployment options to fulfill the business needs and preferences. The following are a few development options:

- **On-Demand Instances:** Without making any advanced commitments, users can easily create the EMR clusters utilizing on demand instances for they need and will pay for the resources on hourly basis. This will be as a flexible choice for shifting workloads well.
- **Reserved Instances:** Reserved Instances are helpful for customers to commit for a specific instance for a duration of 1 or 3 years in a particular region. This option provides an appropriate steady workloads with predictable usage and less expensive than on-demand pricing.
- **Spot Instances:** By using [Amazon EC2 spot instances](#), users can create requests for EC2 capacity that are unused possibly saving a lot of money. Spot instances are best suited for workloads that are tolerant of faults and disrupts.

Advantages Of Amazon EMR

The following are the advantages of amazon EMR:

1. **Scalability:** EMR allows users to easily scale up or down the number of instances in a cluster to handle varying amounts of data processing and analysis tasks.
2. **Cost Effectiveness:** EMR allows users to pay for the resources they need, when they need them, making it a cost-effective solution for big data processing.
3. **Integration With Other AWS Services:** EMR can be easily integrated with other AWS services such as [Amazon S3](#), [Amazon DynamoDB](#), and Amazon Redshift for data storage and analysis.
4. **Flexibility:** EMR supports a wide range of open-source big data frameworks, including Hadoop, Spark, and [Hive](#), giving users the flexibility to choose the tools that best fit their needs.
5. **Easy To Use:** EMR provides an easy-to-use web interface that allows users to launch and manage clusters, as well as monitor and troubleshoot performance issues.

Disadvantages Of Amazon EMR

The following are the disadvantages of Amazon EMR:

1. **Limited Customization:** EMR is pre-configured with popular big data frameworks such as Hadoop and Spark, so users may have limited options for customizing their cluster.
2. **Latency:** The latency of data processing tasks may increase as the size of the data set increases.
3. **Cost:** EMR can be expensive for users with large amounts of data or high-performance requirements, as costs are based on the number of instances and the amount of storage used.
4. **Limited Control Over The Infrastructure:** EMR is a managed service, which means that users have limited control over the underlying infrastructure. This can be a disadvantage for users who need more control over their big data environments.
5. **Limited Support For Certain Big Data Frameworks:** EMR does not support some big data frameworks such as Flink, which may be a deal breaker for some organizations.
6. **Limited Support For Certain Applications:** EMR is not suitable for all types of applications, it mainly supports big data processes and analytics.

Use Cases Of Amazon EMR

The following are the use cases of Amazon EMR:

- **Big Data Processing:** Amazon ECR is ideal for business Organizations where their is a dealing of distributed processing with large amounts of data. It is capable of managing large volumes of data conversions, data [warehousing](#) and analysis of logs efficiently.
- **Data Analysis:** EMR is well known for performing complicated data analytics. It supports with big data frameworks like Apache spark. It facilitates the companies in making well informed decisions by letting them to extract insightful information from various types of datasets.
- **Genomic Analysis:** EMR is used in bio informatics for analyzing genomic data. Large scaled genomic datasets are used for processing and analyzing to helps the researchers in enhancing the scalability and interoperabilities with genomic technologies in life sciences and healthcare.

- **Machine Learning:** EMR supports integration with other AWS services such as Amazon SageMaker seamlessly. It facilitates the organizations to run distributed [ML algorithms](#) on large datasets. Its usage is very beneficial for predictive analysis and model training.

Amazon Redshift

A data warehouse like Amazon Redshift is your best choice when you need to pull together data from many different sources – like inventory systems, financial systems, and retail sales systems – into a common format, and store it for long periods of time. If you want to build sophisticated business reports from historical data, then a data warehouse like Amazon Redshift is the best choice. The query engine in Amazon Redshift has been optimized to perform especially well on running complex queries that join large numbers of very large database tables. When you need to run queries against highly structured data with lots of joins across lots of very large tables, choose Amazon Redshift.

Amazon Redshift is a fast, fully managed data warehousing service in the cloud, enabling businesses to execute complex analytic queries on volumes of data—thus minimizing delays and ensuring sound support for decision-making across organizations. It was released in 2013, built to remedy the problems associated with traditional, on-premises data warehousing, such as scalability, cost, and complexity.

Amazon Redshift is a flexible, massively scalable, cloud-based service that ranges from a few hundred gigabytes of data to several petabytes, it allows businesses to handle increasingly larger data sizes without much upfront investment, the architecture of Redshift is optimized for complex queries and analytics using techniques like columnar storage and massively parallel processing to deliver high-speed query performance.

Amazon Redshift is a fully managed service in the cloud, dealing with petabyte-scale warehouses of data made to store large-scale data and implement effective ways of running even complex queries. Thus, it enables businesses to quickly and cost-effectively analyze huge amounts of data by using SQL-based queries and business intelligence tools.

Key Features of Amazon Redshift

- **Scalability:** Scale from a few hundreds of gigabytes to a petabyte or even more, allowing businesses to grow their data warehouses based on necessity. In its core is a columnar and MPP-based storage that ensures quick query performance, even over large datasets.
- **Integration:** Redshift seamlessly integrates with Amazon S3, Amazon RDS, AWS Glue, and much more to create a data ecosystem.
- **Cost-Effective:** Amazon Redshift is structured in a way such that it turns out to be cost-effective for you, with a couple of pricing options that enable one to pay just for storage and computing power.

How Amazon Redshift Works?

- **Clusters and Nodes:** Redshift groups its resources into clusters. A cluster consists of one or more compute nodes. A leader node manages client connections and SQL processing. Compute nodes execute the queries and store data.
- **Data Storage:** Redshift organizes data in row format followed by organizing it columnar. This architecture minimizes the volume of disk reads and hence increases performance for analytical queries.
- **Query Execution:** Redshift runs each query in parallel on multiple nodes, enabling it to distribute workloads and process large data quantities with MPP architecture.

Use Cases:

- **Business Intelligence:** Companies have large datasets and use Redshift to process complex queries, generate reports, and gain insights into their data for supporting decision-making processes.
- **Data Warehousing:** Primarily, Redshift provides a central data warehouse to store and analyze all data created in various sources.
- **Big Data Analytics:** Since it accommodates petabyte-scale data capacity, Redshift is large enough for an enterprise to analyze big data that allows them to observe any trends or patterns within their data.

What is AWS Glue?

AWS Glue is a serverless data integration service that makes it easy for analytics users to discover, prepare, move, and integrate data from multiple sources. You can use it for analytics, machine learning, and application development. It also includes additional productivity and data ops tooling for authoring, running jobs, and implementing business workflows.

With AWS Glue, you can discover and connect to more than 70 diverse data sources and manage your data in a centralized data catalog. You can visually create, run, and monitor extract, transform, and load (ETL) pipelines to load data into your data lakes. Also, you can immediately search and query cataloged data using Amazon Athena, Amazon EMR, and Amazon Redshift Spectrum.

AWS Glue consolidates major data integration capabilities into a single service. These include data discovery, modern ETL, cleansing, transforming, and centralized cataloging. It's also serverless, which means there's no infrastructure to manage. With flexible support for all workloads like ETL, ELT, and streaming in one service, AWS Glue supports users across various workloads and types of users.

Also, AWS Glue makes it easy to integrate data across your architecture. It integrates with AWS analytics services and Amazon S3 data lakes. AWS Glue has integration interfaces and job-authoring tools that are easy to use for all users, from developers to business users, with tailored solutions for varied technical skill sets.

With the ability to scale on demand, AWS Glue helps you focus on high-value activities that maximize the value of your data. It scales for any data size, and supports all data types and schema variances. To increase agility and optimize costs, AWS Glue provides built-in high availability and pay-as-you-go billing.

AWS Glue features

AWS Glue features fall into three major categories:

- Discover and organize data
- Transform, prepare, and clean data for analysis
- Build and monitor data pipelines

Discover and organize data

- **Unify and search across multiple data stores** – Store, index, and search across multiple data sources and sinks by cataloging all your data in AWS.
- **Automatically discover data** – Use AWS Glue crawlers to automatically infer schema information and integrate it into your AWS Glue Data Catalog.
- **Manage schemas and permissions** – Validate and control access to your databases and tables.
- **Connect to a wide variety of data sources** – Tap into multiple data sources, both on premises and on AWS, using AWS Glue connections to build your data lake.

Transform, prepare, and clean data for analysis

- **Visually transform data with a job canvas interface** – Define your ETL process in the visual job editor and automatically generate the code to extract, transform, and load your data.
- **Build complex ETL pipelines with simple job scheduling** – Invoke AWS Glue jobs on a schedule, on demand, or based on an event.
- **Clean and transform streaming data in transit** – Enable continuous data consumption, and clean and transform it in transit. This makes it available for analysis in seconds in your target data store.
- **Deduplicate and cleanse data with built-in machine learning** – Clean and prepare your data for analysis without becoming a machine learning expert by using the FindMatches feature. This feature deduplicates and finds records that are imperfect matches for each other.
- **Built-in job notebooks** – AWS Glue job notebooks provide serverless notebooks with minimal setup in AWS Glue so you can get started quickly.
- **Edit, debug, and test ETL code** – With AWS Glue interactive sessions, you can interactively explore and prepare data. You can explore, experiment on, and process data interactively using the IDE or notebook of your choice.
- **Define, detect, and remediate sensitive data** – AWS Glue sensitive data detection lets you define, identify, and process sensitive data in your data pipeline and in your data lake.

Build and monitor data pipelines

- **Automatically scale based on workload** – Dynamically scale resources up and down based on workload. This assigns workers to jobs only when needed.
- **Automate jobs with event-based triggers** – Start crawlers or AWS Glue jobs with event-based triggers, and design a chain of dependent jobs and crawlers.
- **Run and monitor jobs** – Run AWS Glue jobs with your choice of engine, Spark or Ray. Monitor them with automated monitoring tools, AWS Glue job run insights, and AWS CloudTrail. Improve your monitoring of Spark-backed jobs with the Apache Spark UI.
- **Define workflows for ETL and integration activities** – Define workflows for ETL and integration activities for multiple crawlers, jobs, and triggers.

Accessing AWS Glue

You can create, view, and manage your AWS Glue jobs using the following interfaces:

- **AWS Glue console** – Provides a web interface for you to create, view, and manage your AWS Glue jobs.
- **AWS Glue Studio** – Provides a graphical interface for you to create and edit your AWS Glue jobs visually.
- **AWS Glue section of the AWS CLI Reference** – Provides AWS CLI commands that you can use with AWS Glue.
- **AWS Glue API** – Provides a complete API reference for developers.

Use Cases of AWS Glue

- **To build Data Warehouse to Organize, Cleanse, Validate, and Format Data:** We can transform and move AWS cloud data into our data store. We can also load data from different sources into our data warehouse for regular reporting and analysis. By storing it in the warehouse, we integrate information from different parts of our business and form a common source of data for decision-making.
- **When we run Serverless Queries against our Amazon S3 Data Link:** S3 here means simple storage service. AWS Glue can catalog our simple storage service that is Amazon S3 data making it available for querying with [Amazon Athena](#) and Amazon RedShift Spectrum. With crawlers, our metadata stays in synchronization with the underlying data. AWS RedShift Spectrum can access and analyze data through one unified interface without loading it into multiple data.
- **Creating event-driven ETL Pipelines:** We can run our ETL jobs as soon as new data becomes available in [Amazon S3](#) by invoking our AWS Glue ETL jobs from an AWS Lambda function. We can also register this new data in the AWS load data catalog as a part of our details.
- **To understand our Data Assets:** We can store our data using various AWS services and still maintain a unique, unified view of our data using the AWS Glue data catalog. We can view the data catalog to quickly search and discover the datasets that we own and maintain the relative data in one central location.

Benefits of AWS Glue


- **Less Hassle:** AWS Glue is integrated across a wide range of AWS services. AWS Glue natively supports data stored in [Amazon Aurora](#) and other [Amazon Relational Database Service](#) engines, Amazon RedShift and Amazon S3 along with common database engines and databases in our virtual private cloud running on [Amazon EC2](#).
- **Cost Effective:** AWS Glue is serverless. There is no infrastructure to provision or manage AWS Glue handles, provisioning, configuration, and scaling of the resources required to run our ETL jobs. We only pay for the resources that we use while our jobs are running.
- **More Power:** AWS Glue automates much of the effort in building, maintaining, and running ETL jobs. It identifies data formats and suggests schemas and transformations. Glue automatically generates the code to execute our data transformations and loading processes.

Disadvantages of AWS Glue

- **Amount of Work Involved:** It is not a full-fledged ETL service. Hence in order to customize the services as per our requirements, we need experienced and skillful candidates. And it involves a huge amount of work to be done as well.
- **Platform Compatibility:** AWS Glue is specifically made for the AWS console and its subsidiaries. And hence it isn't compatible with other technologies.
- **Limited Data Sources:** It only supports limited data sources like S3 and JDBC
- **High Skillset Requirement:** AWS Glue is a serverless application, and it is still a new technology. Hence, the skillset required to implement and operate the AWS Glue is high.

AWS - Amazon Kinesis

Amazon Kinesis is a service provided by Amazon Web Service that allows users to process a large amount of data (which can be audio, video, application logs, website clickstreams, and IoT telemetry) per second in real time. In today's scenario handling a large amount of data becomes very important and for that, there is a complete whole subject known as Big Data which works upon how to process or handle the streams of large amounts of data. So Amazon came up with a solution known as Amazon Kinesis which is fully managed and automated and can handle the real-time large streams of data with ease. It allows users to collect, store, capture, and process many logs from distributed streams such as social media feeds.



AWS Amazon Kinesis Working

Amazon Kinesis working can be divided into **4 stages** as follows.

1. Data Ingestion

Amazon kinesis will collect the data or receives the data from the different data streams like application, sensors, and so on. The data that is going to be received from the different sources can be in different formats like JSON and Binary. It can also accept the data of real-time applications.

2. Sharding and Scaling

The smaller parts of the data called the shards the data which is received from the different sources are divided into smaller shards for redundancy and fault tolerance. There are no limits for the shards amazon kinesis can scale the shards horizontally depending on the requirement.

3. Processing and buffering


After sharded the data will be prepared for further use like it will apply filtering or record aggregation before storing it.

4. Making the data accessible

After completing all the steps mentioned above know the data should be accessible it offers various ways to access and utilize your data stream.

- Kinesis Data Streams API.
- Kinesis Firehose.
- Kinesis Analytics.

Below is the Image that Explain Exactly how Kinesis work in real time:



AWS Kinesis Data Streams

Amazon Kinesis will stream the data in real-time help in handling it and also tell you what to do with that data according to the organization's goals following are the broad categories to get you started:

- **Real time data ingestion and processing:** Amazon kinesis will take the data from the real time it will helps in the applications like health which are used for the identifying the and regulating the health of the patient and also if any any emergency with the help of data we can predict it in before head. It can also used in the applications of OTT platforms by which you can personalize the according to the user experience.
- **Streamlined data delivery and storage:** The real time data can be stored in the storage and can be used for the further research and can be used for the further use and also amazon kinesis can be integrated with the other services also.
- **Real-time insights and automation:** The data which is collected from the real time data will be analyzed the whole data and recats to anomalies, fraud attempts or any other critical immediately. And also monitor the key metrics which can be used for the data driven decision making.

Types of Services Offered by Amazon Kinesis

There are **4 types of services** which Amazon kinesis offers are as follows:

1. Amazon Kinesis Data Streams
2. Amazon Kinesis Video Streams
3. Amazon Kinesis Firehose
4. Amazon Kinesis Data Analytics

1. Amazon Kinesis Data Streams

It provides a platform for real-time and continuous processing of data. It is also used to encrypt the sensitive data by using the KMS master keys and the server-side encryption for the security purpose. The architecture of Amazon Kinesis looks somewhat like the given below image.

2. Amazon Kinesis Video Streams

Amazon Kinesis Video Streams is a fully managed service makes it easy to secure stream video data from devices to the cloud for real-time processing and analytics. With Kinesis Video Streams you can effortlessly capture, store and also process video streams for various use cases including surveillance, media streaming and also IoT applications. It enables seamless integration with machine learning models for video analysis and provides powerful tools for real-time video data processing without worrying about scalability or storage limitations.

Amazon Kinesis video streams is a power tool that is provided by AWS as a service that can deliver live on-demand streams in the real-time following are some of the key features used with kinesis streams.

- Ingest live video from various sources
- Process video in real-time
- Analyze video content
- Deliver video to various destinations
- Build custom video applications

3. Amazon Kinesis Data Analytics

It allows the streams of data provided by the kinesis firehose and kinesis streams to analyze and process it with the standard SQL. It analyzes the data format and automatically parses the data and by using some standard interactive schema editor to edit it in recommend schema. It also provides pre-built stream process templates that can be used to select a suitable template for their data analytics.

4. Amazon Kinesis Firehouse

Firehouse allows the users to load or transformed their streams of data into amazon web service latter transfer for the other functionalities like analyzing or storing. It does not require continuous management as it is fully automated and scales automatically according to the data.

Amazon Kinesis Features

- **Cost-efficient:** All the services provided by the amazon are cost-efficient as it follows the pay as you go model which means you have to pay for the service according to the usage, not a flat price. So it becomes advantageous for the users that they have to pay only what they use.
- **Integrate with other AWS services:** Amazon Kinesis allows users to use the other AWS services and integrate with it. Services that can be integrated are Amazon DynamoDB, Amazon Redshift, and all the other services that deal with the large amount of data.
- **Availability:** You can access it from anywhere and anytime. Just need a good connectivity of net.
- **Real-time processing-** It allows you to work upon the data which is needed to be updated every time with changes instantaneously. Most advantageous feature of Kinesis because real-time processing becomes important when you are dealing with such a huge amount of data.

Amazon Kinesis Use Cases

- **Real-time application monitoring:** Amazon kinesis will provide the real time data of the applications like if you consider the health application it will provides the live feed of the data by which you can take care of the health by which the issues that is pointed by amazon kinesis.
- **Fraud detection and prevention:** Amazon Kinesis will helps you to protect the data from fraudulent activity by analyzing transaction data by which you can detect the suspicious patterns and blocks fraudulent transactions before they happen.
- **Personalized recommendations and marketing:** Amazon kinesis will helps you in analyzing the data of the customers by which you can understand your customers very better. You can recommends the personalised products in real time to the costumers.
- **IoT analytics and predictive maintenance:** Your connected gadgets' full potential is unlocked with Kinesis. Through the examination of sensor data from electronics, automobiles, or machinery.

Amazon Kinesis Limitations

- The limitation that Amazon kinesis has that it only access the stream of records log for 24 hours by default but it can extend but up to only 7 days not longer than that.
- There is no upper limit in the number of streams that can users have in their accounts.
- One shard supports up to 1000 PUT records per second.

AWS Quicksight

AWS Quicksight is one of the most powerful Business Intelligence tools which allows you to create interactive dashboards within minutes to provide business insights into the organizations. There are number of visualizations or graphical formats available in which the dashboards can be created. The dashboards get automatically updated as the data is updated or scheduled. You can also embed the dashboard created in Quicksight to your web application.

With the latest ML insights, also known as Machine Learning insights, Quicksight uses its inbuilt algorithms to find any kind of anomalies or peaks in the historical data. This helps to get prepared with the business requirements ahead of time based on these insights. Here is quick guide to get started with Quicksight.

AWS Quicksight is an AWS based Business Intelligence and visualization tool that is used to visualize data and create stories to provide graphical details of the data. Data is entered as dataset and you can apply filters, hierarchies, and columns to prepare documents. You can choose various charts like Bar charts, Pie charts, etc. to visualize the data effectively. This basic tutorial will help you to understand and learn AWS Quicksight tool.

Amazon QuickSight is a cloud-scale business intelligence (BI) service that you can use to deliver easy-to-understand insights to the people who you work with, wherever they are. Amazon QuickSight connects to your data in the cloud and combines data from many different sources. In a single data dashboard, QuickSight can include AWS data, third-party data, big data, spreadsheet data, SaaS data, B2B data, and more. As a fully managed cloud-based service, Amazon QuickSight provides enterprise-grade security, global availability, and built-in redundancy. It also provides the user-management tools that you need to scale from 10 users to 10,000, all with no infrastructure to deploy or manage.

QuickSight gives decision-makers the opportunity to explore and interpret information in an interactive visual environment. They have secure access to dashboards from any device on your network and from mobile devices.

Why QuickSight?

Every day, the people in your organization make decisions that affect your business. When they have the right information at the right time, they can make the choices that move your company in the right direction.

Here are some of the benefits of using Amazon QuickSight for analytics, data visualization, and reporting:

- The in-memory engine, called SPICE, responds with blazing speed.
- No upfront costs for licenses and a low total cost of ownership (TCO).
- Collaborative analytics with no need to install an application.
- Combine a variety of data into one analysis.
- Publish and share your analysis as a dashboard.
- Control features available in a dashboard.
- No need to manage granular database permissions—dashboard viewers can see only what you share.

For advanced users, QuickSight Enterprise edition offers even more features:

- Saves you time and money with automated and customizable data insights, powered by machine learning (ML). This enables your organization to do the following, without requiring any knowledge of machine learning:
 - Automatically make reliable forecasts.

- Automatically identify outliers.
- Find hidden trends.
- Act on key business drivers.
- Translate data into easy-to-read narratives, like headline tiles for your dashboard.
- Provides extra Enterprise security features, including the following:
 - Federated users, groups, and single sign-on (IAM Identity Center) with AWS Identity and Access Management (IAM) Federation, SAML, OpenID Connect, or AWS Directory Service for Microsoft Active Directory.
 - Granular permissions for AWS data access.
 - Row level security.
 - Highly secure data encryption at rest.
 - Access to AWS data and on-premises data in Amazon Virtual Private Cloud
- Offers pay-per-session pricing for the users that you place in the "reader" security role—*readers* are dashboard subscribers, people who view reports but don't create them.
- Empowers you to make QuickSight part of your own websites and applications by deploying embedded console analytics and dashboard sessions.
- Makes our business your business with multitenancy features for value-added resellers (VARs) of analytical services.
- Enables you to programmatically script dashboard templates that can be transferred to other AWS accounts.
- Simplifies access management and organization with shared and personal folders for analytical assets.
- Enables larger data import quotas for SPICE data ingestion and more frequently scheduled data refreshes.


AWS Database Migration Service (DMS)

The AWS Database Migration Service (DMS) helps you quickly and safely transfer your databases. It also lets you build, analyze, transform, and relocate databases and analytics platforms all in one place, saving you time, resources, and money. In that process, application downtime, based on the source database, is minimal. This may also include Database Migration Service supporting migrations to and from most common commercial and open source databases.

AWS Database Migration Service allows the movement of data across databases. AWS Database Migration Service makes it possible to replicate data constantly. Continuous data replication serves a variety of purposes, including synchronizing Dev/Test environment locales, global database dispersion, and Disaster Recovery instance locations. The service also allows continuous data replication to keep the source and target databases synchronized during the migration process.


AWS DMS - Database Migration Service Working

- It is a web service to help you migrate data from one data store to another within a minimal period.
- You can migrate between source and target endpoints, which are the same kind of database engine.
- The example would be from an Oracle Database to another Oracle Database. The only requirement to use AWS DMS is one of your endpoints must be on this service.
- You may pay for it as you use it. Neither the initial license purchase money nor the recurring maintenance fee are required.



AWS DMS - Database Migration Service Components

Understanding the key components of AWS DMS is crucial for planning and executing a successful data migration. Below is an illustrative diagram showing the components used in an AWS DMS data migration



This image highlights the essential components and how they interact throughout the migration process. Let's break down these components and their roles in detail.

1. Source Database (Source DB)

The **Source Database** is where your original data resides before migration. AWS DMS supports a wide range of database types for the source, including on-premises databases, Amazon Relational Database Service (RDS), or data stored in other AWS services.

Example: You might have data in an on-premises MySQL database that you plan to migrate to AWS.

2. Source Endpoint

A **Source Endpoint** defines the connection between AWS DMS and your source database. It includes the configuration details and authentication credentials required for AWS DMS to access the source database.

Key Responsibilities:

- Ensuring that AWS DMS can securely read data from the source database.
- Managing any necessary conversions or configurations to facilitate data extraction.

3. Replication Instance

The **Replication Instance** is the core engine of AWS DMS. It connects to both the source and target endpoints to migrate data. All data extraction, transformation, and loading tasks run on this instance.

Key Responsibilities:

- Extracting data from the source database.
- Performing any data transformations required.
- Loading data into the target database.

Note: You can choose from various replication instance types and sizes depending on the amount of data and the performance requirements of your migration.

4. Migration Tasks (Replication Tasks)

Migration Tasks, also known as **Replication Tasks**, define how data is moved from the source to the target. Each task controls the data flow and can be customized based on your migration needs.

- **Types of Migration Tasks:**
 - **Full Load Tasks:** Migrates the entire dataset from the source to the target. This is typically used for the initial data load.
 - **Change Data Capture (CDC) Tasks:** Continues to capture changes (inserts, updates, deletes) made to the source database and replicates them to the target in real-time.
- **Customization Options:**
 - Filtering data to migrate only specific schemas or tables.
 - Transforming data (e.g., renaming columns, changing data types) during the migration process.

5. Target Endpoint

A **Target Endpoint** defines the connection to your target database. Like the source endpoint, it includes the settings and credentials required for AWS DMS to write data to the target location.

Key Responsibilities:

- Ensuring secure and successful data loading into the target database.
- Managing configurations that affect how data is written to the target, such as compression or encryption settings.

6. Target Database (Target DB)

The **Target Database** is where the data will reside after the migration. AWS DMS supports a variety of target databases, which can be in the AWS cloud (like Amazon RDS, Amazon Aurora, or Amazon Redshift) or on-premises.

Example: If you are migrating data from an on-premises MySQL database to AWS, the target database might be Amazon RDS for MySQL or Amazon Aurora.

How These Components Interacts?

1. **Source DB and Source Endpoint:** AWS DMS uses the source endpoint to securely connect to and read data from the source database.
2. **Replication Instance and Migration Tasks:** The replication instance hosts migration tasks that carry out the data movement from the source to the target. Depending on the task type, the replication instance either performs a full load, CDC, or both.
3. **Target Endpoint and Target DB:** The migration tasks write the data to the target database through the target endpoint ensuring data integrity and consistency.

AWS Database Migration Service Advantages

Following are the advantages of the AWS Database Migration Service

- **Low cost:** Data Migration Service is a free migration solution for [DocumentDB](#), [Redshift](#), [Aurora](#), and [DynamoDB](#).
- **Serverless:** AWS DMS is serverless, and hence can deploy, maintain, and monitor all hardware and software that it needs for migration, which frees you from traditional tasks such as analyzing capacity, hardware and software procurement, system installation, and administration.
- **Reliability:** Database Migration Service is a self-curing service that immediately restarts in case an interruption occurs.
- **Minimal downtime:** The source remains operational while DMS continuously sends changes throughout the migration process to your data source.

AWS Database Migration Service Limitations

The following are the limitations of the AWS Database Migration Service

- **Performance:** AWS DMS can load up to eight tables in parallel. This can be increased somewhat by using a large replication server.

- **Engine Name:** The name of the target engine that the Fleet Advisor should use in its recommendation.
- **Data migration:** AWS DMS provides security for your data migration. As the data moves from source to target, you can also encrypt it in flight during migration using [Secure Socket Layers \(SSL\)](#).
- **Data gathering:** Change data. Capturing incremental loads in AWS DMS requires some programming, therefore, it takes more time and is labor-intensive.

AWS Server Migration Service

AWS Server Migration Service (SMS) is a cloud service provided by Amazon Web Services (AWS) that enables you to migrate on-premises servers, virtual machines (VMs), and applications to the AWS Cloud. It helps streamline the process of migrating your existing workloads, reducing the complexity and downtime associated with traditional migration methods.

Key features and benefits of AWS Server Migration Service:

1. **Server and VM Migration:** AWS SMS supports migration of both physical servers and VMs from various sources, including VMware vSphere, Microsoft Hyper-V, and Microsoft Azure. It allows you to consolidate and migrate these servers to the AWS Cloud.
2. **Continuous Replication:** AWS SMS offers continuous replication, which means your source servers' data is kept up-to-date in the AWS Cloud, reducing the risk of data loss during the migration process.
3. **Incremental Replication:** After the initial replication, AWS SMS only replicates the changes made to the source servers, minimizing the amount of data transferred and reducing migration time.
4. **Automated Replication and Testing:** The service automates the replication and validation of your server settings, making sure that they are compatible with the AWS environment before the migration.
5. **Flexible Cutover:** You can schedule cutover for your migration, allowing you to choose the right time to switch your applications from the source to the target environment.
6. **Support for Large-scale Migrations:** AWS SMS is designed to handle large-scale migrations, making it suitable for enterprises with extensive on-premises infrastructure.
7. **Integration with AWS Application Discovery Service:** AWS SMS can work in conjunction with AWS Application Discovery Service to provide a better understanding of your on-premises applications, dependencies, and resource utilization before planning the migration.

It is essential to plan your migration carefully, considering factors like network connectivity, security, and application dependencies to ensure a successful migration to the AWS Cloud. AWS SMS can simplify the process and help you achieve a smooth and efficient migration while minimizing downtime and risk.

What is AWS Snowball

AWS Snowball is a petabyte-scale offline data transfer device designed to securely and quickly migrate large datasets to and from Amazon Web Services (AWS). Instead of transferring data over the internet, Snowball allows organizations to physically ship encrypted storage devices for efficient data ingestion.

AWS Snowball Variants

AWS offers a couple of different Snowball models to fit needs:

- **Snowball Edge Storage Optimized** – Designed for large-scale **data migration and storage**, offering **up to 80 TB of usable storage per device**.
- **Snowball Edge Compute Optimized** – Includes **built-in compute capabilities** such as **AWS Lambda and EC2 instances** for **local data processing** before cloud transfer.

These **rugged, portable devices** are ideal for industries dealing with **remote storage, edge computing, and massive offline data transfers**.

Features of Amazon Snowball

Amazon Snowball is part of the AWS Snow Family, which includes other devices like Snowball Edge, Snowcone, and Snowmobile. It is rugged and secure and physical device designed to help businesses move large amounts of data into AWS Quickly, without relying on slow internet connections. This makes it perfect for large-scale data transfers, disaster recovery, media archives and edge computing.

Here are the key features that makes AWS Snowball a standout solution:

1. High-Speed Offline Data Transfer

AWS Snowball transfers **terabytes or petabytes of data** up to **10 times faster** than traditional internet-based methods. It uses **rugged, tamper-proof storage** for secure and reliable transport.

2. Secure and Encrypted Storage

Data is secured with **256-bit encryption** through **AWS Key Management Service (KMS)**. Snowball's **tamper-proof hardware** ensures the security of your data during offline migration.

3. Built-In Compute for Edge Processing

The **Snowball Edge Compute Optimized** version includes support for **AWS Lambda** and **EC2 instances**. This allows you to **process data locally** before migrating it to the cloud—ideal for remote locations and edge computing applications.

4. Easy Integration with NFS Support

AWS Snowball functions as an **NFS (Network File System) mount point**, making it easy to integrate with your on-premises servers and applications. This helps streamline **file-based data migration** to AWS while preserving file system metadata.

5. S3 Compatibility

AWS Snowball is **S3-compatible**, ensuring your applications can interact with your data using standard **S3 APIs**. This integration allows for seamless data management and smooth workflows.

6. GPU Support for High-Performance Workloads

The device supports **GPU capabilities** for processing complex tasks like **machine learning**, **video analysis**, and other high-performance computing applications. This is perfect for resource-intensive tasks in remote environments.

7. Automatic Data Erasure

Once the data is successfully uploaded to AWS, the device automatically **erases all data** from the Snowball, ensuring that no sensitive information is left behind. This provides **complete protection** for your data.

8. Automatic E Ink Shipping Label

The Snowball device comes with an **E Ink shipping label** that automatically updates with the correct destination address, reducing the chances of logistics errors and simplifying the return process.

9. Cost-Efficient

AWS Snowball offers a **cost-effective solution** for large-scale data movement. It avoids the **expensive fees** typically associated with high-bandwidth internet-based transfers, making it an affordable alternative for businesses.

10. Flexible Online & Offline Transfer Options

- Offline: Ship Snowball devices for bulk data transfer without an internet connection.
- Online: Use AWS DataSync for real-time hybrid cloud synchronization.

11. Clustering for Scalable Storage

You can **cluster multiple Snowball devices** to increase both **storage capacity** and **durability**. This ensures resilient data storage solutions for **edge computing**, even if one device fails.

How AWS Snowball Works

AWS Snowball simplifies **large-scale data migration** by allowing you to transfer data securely using a **physical device** instead of relying on slow internet connections. Follow the below steps to use AWS Snowball:

Step 1: Order the Snowball Device

- Log in to the [AWS Snowball Console](#).
- Choose the type of device:
 - **AWS Snowball Edge Compute Optimized** (for running applications at the edge).
 - **AWS Snowball Edge Storage Optimized** (for large-scale data transfer).
- Create a job:
 - Select an **Amazon S3 bucket** where your data will be stored.
 - Enable [Amazon Simple Notification Service \(SNS\)](#) for tracking updates.
 - Configure options such as **Amazon EC2 AMIs** if needed.

- **AWS Prepares and Ships the Device:**
 - AWS configures the device and **ships it to your location**.

Step 2: Set Up the Device

- **Unpack the AWS Snowball** upon arrival.
- **Power it up** and connect it to your **local area network (LAN)**.
- **Use AWS OpsHub** (a graphical user interface) to:
 - Unlock the device securely.
 - Manage the device settings.
 - Start the data transfer or launch EC2 instances if using the compute-optimized version.

Step 3: Transfer Your Data

- **Use AWS OpsHub** or the **AWS Snowball Client** to move data from your local storage to the Snowball device.
- Monitor the transfer process to ensure all files are securely copied.

Step 4: Return the Device to AWS

- Once the data transfer is complete, **shut down the device**.
- **Prepare it for return shipping:**
 - The **E Ink shipping label automatically updates** to display the correct return address.
- **Ship the device back to AWS** using the pre-paid label.

Step 5: Data Upload & Secure Erasure

- **AWS receives the Snowball device** at its data center.
- Your data is **automatically transferred** from the Snowball device to your **Amazon S3 bucket**.
- AWS **verifies the data transfer** to ensure completeness.
- Once the transfer is complete, AWS **securely erases** all data from the device and **sanitizes it**, ensuring no customer data remains.

Real-World Use Cases for AWS Snowball

AWS Snowball is ideal for a variety of industries and use cases. Here are some examples:

- **Cloud Migration:** If your company needs to move large datasets from on-premises storage to the cloud, Snowball makes it fast and easy.
- **Disaster Recovery:** Secure backups of business-critical data can be sent to AWS Snowball for long-term storage and easy recovery.

- **Media & Entertainment:** Snowball is commonly used by movie studios, game developers, and media companies to transfer **high-resolution videos** and raw media files for cloud processing.
- **Edge Computing:** Industries like **oil and gas**, **mining**, and **research** use Snowball Edge to process data locally before transferring it to the cloud.
- **Healthcare & Genomics:** Medical institutions and research facilities use Snowball to move **large genomic datasets** securely and compliantly to AWS for analysis.

Benefits of AWS Snowball

1. Simple Migration

AWS Snowball makes moving large amounts of data easy. With a pre-configured device, you can load your data, ship it back to AWS, and let the service handle the transfer, saving time and avoiding complex setups.

2. Faster Performance

Snowball accelerates data migration, allowing you to move data much faster than over the internet. It reduces transfer time, making large-scale migrations more efficient.

3. Strong Data Protection

Your data is safe with AWS Snowball. It's encrypted with 256-bit encryption, and once the device is returned, the data is securely erased, ensuring full protection throughout the process.

AWS Snowball Vs. Other AWS Data Transfer Services

AWS provides multiple data transfer solutions to meet different business needs. Below is a comparison of AWS Snowball, AWS DataSync, AWS Snowmobile, and AWS Direct Connect to help you choose the best option.

Feature	AWS Snowball	AWS DataSync	AWS Snowmobile	AWS Direct Connect
Best For	Offline bulk data transfer	Automated online sync	Exabyte-scale migration	Dedicated high-speed link
Transfer Speed	Fast (Days for TBs-PBs)	Medium (Depends on bandwidth)	Ultra-fast (Truck-based)	Up to 100 Gbps
Capacity	Up to 80 TB per device	Unlimited (Depends on internet speed)	Up to 100 PB per truck	Up to 100 Gbps per connection

Feature	AWS Snowball	AWS DataSync	AWS Snowmobile	AWS Direct Connect
Security	256-bit encryption, tamper-proof	End-to-end encryption	Highly secure, surveillance-equipped	Private, encrypted network
Internet Needed?	No	Yes	No	Yes
Setup Complexity	Simple (Order, transfer, return)	Moderate (Requires configuration)	Complex (Requires AWS approval)	Complex (Needs network setup)
Use Cases	Cloud migration, disaster recovery, remote locations	Regular file transfers, hybrid cloud sync	Massive enterprise migrations	Real-time data transfer, low-latency workloads
Cost Efficiency	Avoids high network costs	Varies by data volume	Expensive, used for extreme cases	Best for long-term, large-scale use

Features of Snowball

AWS Snowball with the snowball device has the following features:-

- You don't need to purchase or maintain your own hardware.
- Your data is secure both while it's at rest and when it's being physically moved.
- You can manage your jobs programmatically using the job management API or through the AWS Snow Family Management Console.
- The 50 TB model is accessible in all other AWS Regions; 80 TB and 50 TB variants are only available in US Regions.
- Local data transfers are possible between a Snowball and an on-premises data center. These transfers can be carried out using the standalone, downloadable Snowball client. Or you can use the free Amazon S3 Adapter for Snowball to transfer data programmatically using calls to the Amazon S3 REST API. See Data Transmission with a Snowball for further details.


- When the Snowball is prepared to ship, the E Ink display on its own shipping container switches to display your shipping label. See Shipping Factors for AWS Snowball for more details.

What is AWS Snowmobile?

AWS Snowmobile is a physical data transfer service designed for large-scale data migrations to the AWS cloud. It is essentially to secure 45-foot long shipping container that can transfer up to **100 petabytes** of data in a single shipment. AWS Snowmobile is ideal for industries needing to move vast amounts of data such as media, entertainment or government sectors. It offers high-level security including encryption GPS tracking and security personnel to ensure safe transport of your data to AWS data centers.

How does AWS Snowmobile Work?

AWS Snowmobile works by physically transporting large amounts of data from your location to AWS data centers. Refer to the diagram below for a visual explanation On How does AWS Snowmobile Work.



Use Cases For AWS Snowmobile

AWS Snowmobile is ideal for large-scale data migrations where transferring massive amounts of data quickly is crucial. It is commonly used in the following scenarios:

- **Data center migration:** When organizations need to move petabytes or exabytes of data to the cloud Snowmobile offers a fast and secure solution.
- **Disaster recovery:** Companies use Snowmobile to back up vast amounts of data to the cloud ensuring quick recovery in case of a disaster.
- **Video libraries:** Media companies transferring large video archives to AWS for storage and processing benefit from Snowmobile's massive capacity.

- **Research data transfer:** Industries like healthcare and genomics use Snowmobile to transport enormous datasets, enabling advanced cloud analytics and processing.
- **Government and security agencies:** For secure, large-scale data transfer, Snowmobile provides enhanced security features for sensitive data migration.

Benefits of AWS Snowmobile

Benefits of AWS Snowmobile

AWS Snowmobile offers several advantages for organizations with massive data migration needs:

- **Massive Data Transfer:** Snowmobile can transfer exabytes of data making it ideal for large-scale migrations that would otherwise take months or years over standard internet connections.
- **Cost-Effective:** It reduces the time and expense associated with transferring vast amounts of data by eliminating the need for long-term network setups.
- **High Security:** Snowmobile provides end-to-end encryption secure tamper-proof containers and GPS tracking ensuring your data is safe throughout the transfer process.
- **Fast Migration:** By moving data physically rather than over the internet Snowmobile drastically shortens the data migration timeline.
- **Scalability:** Snowmobile is designed to handle the largest data migrations offering scalability for organizations moving entire data centers to the [AWS cloud](#).

Snowball: This is a petabyte-scale data migration device designed to transfer large amounts of data to and from AWS. It's a rugged, secure appliance that can be physically shipped to your location.

Snowball Edge: This device combines compute and storage capabilities, making it suitable for hybrid and edge workloads. It can process data locally and store it before transferring it to AWS.

Snowmobile: For exabyte-scale data migration, AWS offers Snowmobile. It's a semi-trailer truck equipped with petabytes of storage capacity, allowing you to transfer massive amounts of data to and from AWS.



What is Snowball Edge?

Snowball Edge is a device with on-board storage and compute power for select AWS capabilities. Snowball Edge can process data locally, run edge-computing workloads, and transfer data to or from the AWS Cloud.

Each Snowball Edge device can transport data at speeds faster than the internet. This transport is done by shipping the data in the devices through a regional carrier. The appliances are rugged, complete with E Ink shipping labels.

Snowball Edge devices have two options for device configurations—*Storage Optimized 210 TB* and *Compute Optimized*. When this guide refers to Snowball Edge devices, it's referring to all options of the device. When specific information applies only to one or more optional configurations of devices, it is called out specifically

Snowball Edge features

Snowball Edge devices have the following features:

- Large amounts of storage capacity or compute functionality for devices. This depends on the options you choose when you create your job.
- Network adapters with transfer speeds of up to 100 Gbit/second.
- Encryption is enforced, protecting your data at rest and in physical transit.
- You can import or export data between your local environments and Amazon S3, and physically transport the data with one or more devices without using the internet.
- Snowball Edge devices are their own rugged box. The built-in E Ink display changes to show your shipping label when the device is ready to ship.
- Snowball Edge devices come with an on-board LCD display that can be used to manage network connections and get service status information.
- You can cluster Snowball Edge devices for local storage and compute jobs to achieve data durability across 3 to 16 devices and locally grow or shrink storage on demand.
- You can use Amazon EKS Anywhere on Snowball Edge devices for Kubernetes workloads.
- Snowball Edge devices have Amazon S3 and Amazon EC2 compatible endpoints available, enabling programmatic use cases.
- Snowball Edge devices support the new sbe1, sbe-c, and sbe-g instance types, which you can use to run compute instances on the device using Amazon Machine Images (AMIs).
- Snowball Edge supports these data transfer protocols for data migration:
 - NFSv3
 - NFSv4
 - NFSv4.1
 - Amazon S3 over HTTP or HTTPS (via API compatible with AWS CLI version 1.16.14 and earlier)

Services related to Snowball Edge

You can use an AWS Snowball Edge device with the following related AWS services:

- **Amazon S3 adapter** — Use for programmatic data transfer in to and out of AWS using the Amazon S3 API for Snowball Edge, which supports a subset of Amazon S3 API operations. In this role, data is transferred to the Snow device by AWS on your behalf and the device is shipped to you (for an export job), or AWS ships an empty Snow device to you and you transfer data from your on-premises sources to the device and ship it back to AWS (for an import job)"
- **Amazon S3 compatible storage on Snowball Edge** — Use to support the data needs of compute services such as Amazon EC2, Amazon EKS Anywhere on Snow, and others. This feature is available on Snowball Edge devices and provides an expanded Amazon S3 API set and features such as increased resiliency with flexible cluster setup for 3 to 16 nodes, local bucket management, and local notifications.
- **Amazon EC2** – Run compute instances on a Snowball Edge device using the Amazon EC2 compatible endpoint, which supports a subset of the Amazon EC2 API operations.
- **Amazon EKS Anywhere on Snow** – Create and operate Kubernetes clusters on Snowball Edge devices.
- **AWS Lambda powered by AWS IoT Greengrass** – Invoke Lambda functions based on Amazon S3 compatible storage on Snowball Edge storage actions made on an AWS Snowball Edge device.
- **Amazon Elastic Block Store (Amazon EBS)** – Provide block-level storage volumes for use with EC2-compatible instances.
- **AWS Identity and Access Management (IAM)** – Use this service to securely control access to AWS resources.
- **AWS Security Token Service (AWS STS)** – Request temporary, limited-privilege credentials for IAM users or for users that you authenticate (federated users).
- **Amazon EC2 Systems Manager** – Use this service to view and control your infrastructure on AWS.

AWS monitoring of Snowball Edge

- AWS will monitor the Snow device and may collect metrics and usage information when the Snow device is connected to an AWS Region. If the Snow device is not connected to the AWS Region, then AWS will not monitor the Snow device.
- If AWS detects an irreparable issue, and there is a need to replace physical equipment, AWS will notify you. You can then place a replacement job that we will ship to your site. There is no additional charge for this, as Snow device monitoring is included as part of the Snow device service fee.

